

0.1 Portfolio Contents


1. [Excerpt from Team's Portfolio for Judging](#)

The team submits a 15-page portfolio to judges for awards at every competition. This is an excerpt from ours, and is the page that contains the overview of my motion code for the judges. I wrote the movement section of this excerpt.

2. [Explanation of the Math and algorithm](#)

This is a document I prepared that gives a detailed explanation of the algorithms and math that I'm using for the code.

A copy of the source code that we run on the robot can be found [here](#). My motion planning code is in the Motion subdirectory.




AUTONOMOUS OBJECTIVES

We have designed our robot to be as versatile as possible in auto. Because of our 6 wheel drivetrain and rapid-fire shooting, we are capable of scoring 15 artifacts in auto, with the robot starting from either side of the field. A huge part of why this is possible is because of our new movement system (see below), and our localization which filters together data from our camera and our odometry pods, resulting in accurate paths and accurate shooting, which is a must to have autos that can be consistent on top of being high-scoring.

MOVEMENT

As part of our summer research, we developed a brand-new pathing system. Our motion control helped us win control at the 2025 World Championship, but there were significant flaws we wanted to address. Specifically, we wanted to **improve accuracy** and follow **nonlinear trajectories**. Our new system uses a physics-based model of the robot dynamics (right) to iteratively minimize a cost functional and compute **provably time-optimal trajectories**. Using a **Taylor Expansion**, we also compute **rotating and time-varying PDFL** controllers to follow the trajectory and correct random error. We use a variant of **Differential Dynamic Programming** inspired by recent research from the Universities of CalTech, Oxford, Edinburgh, and others. Our programmers enjoyed the advanced math (tensor analysis and multivariate calculus), which enables **faster, more accurate, and flexible** control. (Come to our pit to learn more about the math—these pages are too small to contain it.) Though related algorithms (e.g. iLQR) have been discussed on the FTC and FRC Discord servers, **we are not aware of any other FTC or FRC team that has used this specific technique in competition.**

Predicted and Measured Velocity over Time



This is based purely on our model, with **no feedback control**. Positional accuracy to the target was within 5%, improved to **0.3%** with correction from sensors.

163.2 kB of PDFL parameters, optimally auto tuned with our advanced model:

$$\dot{x}^a = h_b^a(x) (A_c^b x^c + B_i^b s^i(u))$$

Heading transform Current State Control Response Control Input

State change (derivative) Uncontrolled Robot Dynamics Control Limiter

Comparison of FTC Pathing

Road-Runner	Pedro Pathing	WPM 24-25	WPM 25-26
Some	Approximate	Some	Above dynamics
Limited		In theory	Full Support
Recomputes Paths	Always	Reliant	When Needed
Nonlinear Motion	Splines		Shortest-time path
Unique		Like others	First in FIRST
Optimal			Provably

Figure 1: Section of the team's Portfolio used in Judging

Pathing Math

Kieran Mattingly, 18438 Wolfpack Machina

1 Introduction

The FTC game this year is a shooter game, where we have to intake wiffle balls from the field, sometimes autonomously. Our team originally decided to use a differential drivetrain, like a tank. As we can only move forwards/backwards and turn, a good pathing system is necessary to handle fast, precise movement along curved trajectories. For this design, five variables are needed to represent the robot state at a given time: x and y position, θ , the robot heading, v , the one-dimensional velocity along the heading, and the angular velocity $\dot{\theta}$. We have two control parameters that correspond to the PWM duty cycle sent to the left and right motors.

2 The Dynamical Model

First, the robot dynamics are specified as a Partial Differential Equation (PDE). This equation, in index notation using Einstein summation convention, is

$$\dot{x}^a = h_b^a(x)(A_c^b x^c + B_i^b s^i(u))$$

where x^a is the state vector, consisting of $\langle x, y, \theta, v, \dot{\theta} \rangle$. As explained in more detail later, $h_b^a(x)$, A_c^b , B_i^b , $s^i(u)$, and u , are the heading transformation, control-free state evolution, control response, control limiter, and control input (index has been suppressed for notational convenience), respectively. We computed these dynamics based off an (approximate) physical model of DC motors. Because mass and moment of inertia are approximately constant during play, the torque of the motors is directly proportional to the translational and rotational acceleration of the robot chassis. In theory, the proportionality constants could be calculated by weighing the robot, measuring its moment of inertia, and looking up the motor torque, however in practice we find it simpler to tune this manually. Thus, by setting the torques on the wheels (or a value directly proportional to it), we can directly compute the

acceleration.

The other key component about the robot dynamics we've modeled is back EMF. This is caused by the magnetic field inside the DC motor inducing a counter-torque, so at higher velocities the motor has a lower maximum torque. For example, see Figure 1.




Figure 1: A motor curve graph for one of the motors we plan on using. The yellow line represents the motor torque vs. speed. Source: Game Manual 0. (<https://gm0.org/ro/latest/docs/power-and-electronics/motor-guide/motor-power.html>)

Notice that it follows a decreasing linear trend. We modeled this as Stokes' drag, where the drag force is directly proportional to the velocity. However, the torque due to the control input remains constant. Thus, at high velocities, the drag force balances out the motor power, resulting in net-zero torque, and at low velocity the motor has no drag and max torque, exactly the relationship witnessed in the yellow line of the graph. Friction is handled outside this model.

Mathematically, this effect is in the A and B matrices, defined as

$$A = \begin{bmatrix} 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & B_d & 0 \\ 0 & 0 & 0 & 0 & B_\theta \end{bmatrix}, \quad B = \begin{bmatrix} 0 & 0 \\ 0 & 0 \\ 0 & 0 \\ \tau_d & \tau_d \\ -\tau_\theta & \tau_\theta \end{bmatrix}$$

Where τ_d and τ_θ are the constant motor torques on the driving (forward) position and angle. B_d and B_θ represent the negative drag coefficient on velocity in the driving direction and on the angular velocity. In the A matrix is also the definition of the velocity state variables acting on position, in the first two rows. As the velocity is always in the heading direction (θ), the velocity in x and y depends on the heading. To handle this, the A matrix is defined for when $\theta = 0$, and the function $h(x)$ handles rotating the velocity to be in the direction of movement. $h(x)$ is defined to be:


$$h(x) = \begin{bmatrix} \cos(\theta) & -\sin(\theta) & 0 & 0 & 0 \\ \sin(\theta) & \cos(\theta) & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 1 \end{bmatrix}$$

To predict the future state, We approximate the integral of our robot dynamics equation, obtaining a state transition function $f^a(x, u) = x_{t+\Delta t}^a \approx x^a + \dot{x}^a \Delta t = x^a + h_b^a(x)(A_c^b x^c + B_i^b s^i(u))\Delta t$. The $s^i(u)$ function is present to enforce control limits. There's a limit to the power we can send to the motors, so $s^i(u)$ is a sigmoid function (in this case, hyperbolic tangent) to smoothly map from $\mathbb{R} \rightarrow (-1, 1)$ (or whatever we want the control range to be). This is necessary as the control algorithm itself will assume the control input is unbounded, which results in infeasible trajectories. This concludes our dynamic model. To evaluate the accuracy, we generate a path to move forward 100 centimeters, and run it with no feedback control. Without spending too much time tuning (it was only a testing chassis), it was able to get withing 5cm of the target. In our testing, there was too much randomness (even with identical control inputs) to reliably get any closer. With feedback enabled, this shrunk to less than 3mm (Figure 2).

3 The Cost Function

Paths are defined as a two-tuple (X, U) , where X and U are the state and control trajectories, sampled Δt apart. In order to define what makes a

Figure 2: Velocity vs. Time with and Without Feedback Control



(a) Measured (Red) vs Predicted (Blue) velocity over time without feedback control. Final positional accuracy was within 5%. Notice that towards the end the decreased drag force causes the model to smoothly asymptote, and the data confirms this model's accuracy.

(b) Measured (Red) vs Predicted (Blue) velocity over time with feedback control. Notice the increased noise due to the controller correcting small errors, including errors perpendicular to the direction of motion. Final positional accuracy was within 0.3%.

trajectory "better" or "worse," however, a functional $J : X \times U \rightarrow \mathbb{R}$ must be defined. The traditional way this is done for a sequence of n elements is the quadratic cost function:

$$J(X, U) = \ell_f(x) + \sum_{k=1}^n \ell(x, u),$$

where

$$\ell_f(x) = (x_n - x_{target})^T Q_f (x_n - x_{target}),$$

$$\ell(x, u) = (x_k - x_{target})^T Q (x_k - x_{target}).$$

and Q is the ongoing error cost matrix, R is the control cost matrix, and Q_f is the final cost matrix. This is what we currently use (although we do have ideas on how to alter this function to add obstacle avoidance and other features). Intuitively, this function makes sense; it punishes the robot for being far away from the target state for an extended period of time, but doesn't push it to the maximum control. This last part is important because it prevents the trajectory generation from resulting in a "Bang-Bang" controller, where the controls are set to their max power then minimum power. While this would be the theoretical fastest, it is not the best trajectory

as it leaves no room for correction if there is an error in the model or, in this case, uncertainty. By adding in the control cost term R , the controller is discouraged from this behavior, helping to ensure controllability.

4 Path Generation and Derivatives

The primary component of this algorithm is the generation of paths. This is done via a technique called **Differential Dynamic Programming**, developed by Dr. Richard Bellman in the 1950s. The key idea is to, at each time step, consider what change would be the most beneficial to the path, not just at the current state, but in the future. We can recursively construct a *Value* function V for the current time step, by saying that

$$V_k(x) = \min_u [\ell(x, u) + V_{k+1}f(x, u)]$$

where $f(x, u)$ is the state transition function. In effect, this means that the improvement at the current time step is how much it helps at that time step and how much it helps at the next time step. No recursion is complete without a base case, which is defined to be $V_n(x) = \ell_f(x)$, as the final cost function is itself a measure of the cost of being off at the final state. (This is a version of the *Bellman equation*, restated in this particular view.)

To now use this equation, we initialize a nominal stationary path, and repeatedly compute an update iteratively starting from x_n , using Differential Dynamic Programming. Briefly stated, another function Q can be defined to be the quality of an update in both state and control, taking into account robot dynamics. This can be approximated with a second-order Taylor expansion, with the following derivatives:

$$\begin{aligned} \frac{\partial Q}{\partial x^i} &= \frac{\partial \ell}{\partial x^i} + \frac{\partial f^a}{\partial x^i} \frac{\partial V}{\partial x^a} \\ \frac{\partial Q}{\partial u^k} &= \frac{\partial \ell}{\partial u^k} + \frac{\partial f^a}{\partial u^k} \frac{\partial V}{\partial x^a} \\ \frac{\partial^2 Q}{\partial x^i \partial x^j} &= \frac{\partial^2 \ell}{\partial x^i \partial x^j} + \frac{\partial f^a}{\partial x^i} \frac{\partial^2 V}{\partial x^a \partial x^b} \frac{\partial f^b}{\partial x^j} + \frac{\partial V}{\partial x^c} \frac{\partial f^c}{\partial x^i \partial x^j} \\ \frac{\partial^2 Q}{\partial u^k \partial u^l} &= \frac{\partial^2 \ell}{\partial u^k \partial u^l} + \frac{\partial f^a}{\partial u^k} \frac{\partial^2 V}{\partial x^a \partial x^b} \frac{\partial f^b}{\partial u^l} + \frac{\partial V}{\partial x^c} \frac{\partial f^c}{\partial u^k \partial u^l} \\ \frac{\partial^2 Q}{\partial u^k \partial x^j} &= \frac{\partial^2 \ell}{\partial u^k \partial x^j} + \frac{\partial f^a}{\partial u^k} \frac{\partial^2 V}{\partial x^a \partial x^b} \frac{\partial f^b}{\partial x^j} + \frac{\partial V}{\partial x^c} \frac{\partial f^c}{\partial u^k \partial x^j} \end{aligned}$$

Now, the reason for using index notation and Einstein's summation convention becomes clear: there are three tensor multiplication terms, at the final

term in the second derivatives of Q .

From these derivatives, we can compute an optimal update for the control,

$$\delta u^{l*} = - \left(\frac{\partial^2 Q}{\partial u^k \partial u^l} \right)^{-1} \left(\frac{\partial Q}{\partial u^k} + \frac{\partial^2 Q}{\partial u^k \partial x^j} \delta x^j \right)$$

which can be split into a feed-back term and a feed-forward term. At this stage we also calculate the derivatives of our value function,

$$\begin{aligned} \frac{\partial V}{\partial x^i} &= \frac{\partial Q}{\partial x^i} - \frac{\partial^2 Q}{\partial x^j \partial u^k} \left(\frac{\partial^2 Q}{\partial u^k \partial u^l} \right)^{-1} \frac{\partial Q}{\partial u^l} \\ \frac{\partial V}{\partial x^i \partial x^j} &= \frac{\partial^2 Q}{\partial x^i \partial x^j} - \frac{\partial^2 Q}{\partial x^i \partial u^k} \left(\frac{\partial^2 Q}{\partial u^k \partial u^l} \right)^{-1} \frac{\partial^2 Q}{\partial u^l \partial x^j} \end{aligned}$$

In parts of the above math, you see the derivatives of our $\ell(x)$ and $f(x, u)$ functions appear, and while past work has calculated them using finite differences, we computed them all analytically from our model, shown below:

$$\begin{aligned} f^a(x, u) &\approx x^a + h_b^a(x)(A_c^b x^c + B_i^b s^i(u)) \Delta t \\ \frac{\partial f^a}{\partial x^j} &= \delta_j^a + \left(\frac{\partial h_b^a}{\partial x^j} (A_c^b x^c + B_i^b s^i(u)) + h_b^a(x) A_c^b \delta_j^c \right) \Delta t \\ \frac{\partial^2 f^a}{\partial x^j \partial u^k} &= \frac{\partial h_b^a}{\partial x^j} \left(B_i^b \frac{\partial s^i}{\partial u^k} \right) \Delta t \\ \frac{\partial f^a}{\partial u^k} &= h_b^a(x) \left(B_i^b \frac{\partial s^i}{\partial u^k} \right) \Delta t \\ \frac{\partial^2 f^a}{\partial x^j \partial x^k} &= \left(\frac{\partial^2 h_b^a}{\partial x^j \partial x^k} (A_c^b x^c + B_i^b s^i(u)) + \frac{\partial h_b^a}{\partial x^j} A_c^b \delta_k^c + \frac{\partial h_b^a}{\partial x^k} A_c^b \delta_j^c \right) \Delta t \\ \frac{\partial^2 f^a}{\partial u^j \partial u^k} &= h_b^a(x) \left(B_i^b \frac{\partial^2 s^i}{\partial u^j \partial u^k} \right) \Delta t \end{aligned}$$

We did the same with our cost function:

$$\begin{aligned} \ell(x, u) &= (t_a - x_a) Q_b^a (t^b - x^b) + u_c R_d^c u^d \\ \frac{\partial \ell}{\partial x^i} &= -\delta_{ai} Q_b^a (t^b - x^b) - (t_a - x_a) Q_b^a \delta_i^b \\ \frac{\partial \ell}{\partial u^i} &= \delta_{ci} R_d^c u^d + u_c R_d^c \delta_i^d \\ \frac{\partial \ell}{\partial x^i \partial x^j} &= \delta_{ai} Q_b^a \delta_j^b + \delta_{aj} Q_b^a \delta_i^b \\ \frac{\partial \ell}{\partial u^i \partial u^j} &= \delta_{ci} R_d^c \delta_j^d + \delta_{cj} R_d^c \delta_i^d \\ \frac{\partial \ell}{\partial u^i \partial x^j} &= 0 \end{aligned}$$

Now, we can simulate forward the control sequence, and get a new path. We initialize $\hat{x}_1 = x_1$, and say $\hat{u}_k = u_k - \left(\frac{\partial^2 Q}{\partial u^k \partial u^l}\right)^{-1} \left(\frac{\partial Q}{\partial u^k} + \frac{\partial^2 Q}{\partial u^k \partial x^j}(\hat{x}_k - x_k)\right)$ per our rule above. We populate $\hat{x}$ with the system dynamics, $\hat{x}_k = f(x_{k-1} + u_{k-1})$. After several iterations of the backwards and forwards path, the trajectory reaches a minimal state. To improve convergence, however, a regularization term is added to the eigenvalues of the $\frac{\partial^2 Q}{\partial u^k \partial u^l}$ Jacobian. When the update reduces the cost, the regularization term is decreased, as the approximation is closer to a true quadratic, and if it is not a better trajectory then the regularization term is increased. This is called a *Levenberg-Marquardt* heuristic, and interpolates between using information from the second derivative, or ignoring it and only trusting the first derivative.

While following the trajectory, it is not a good idea to recompute the optimal path, as we want our robot to be consistent, following the same trajectory every time, and we are also limited in processing power. Thus, we instead store the control update function, $\delta u^{l*}(\delta x^a)$, and use that to generate the feedback controller based off the robot's state error. This function is, in effect, a Taylor expansion of our optimal control input around our trajectory at every timestep. As our state contains both position and it's derivative, velocity, it can be thought of as a PD controller, which combined with our model feedforward and friction loopback correction, forms a time-varying PDFL controller. Given that we need to call this function at any time, not just at the set timestamps, we linearly interpolate between the nearest (in time) two points on the trajectory.

However, due to our highly nonlinear model, the higher-order terms of the Taylor expansion have a sizeable effect. If we just used the first-order expansion, then consider what the behavior would be when the robot was 180° offset from the target path (a particularly bad, but illustrating, example). Then, any error in the drivetrain direction would be compensated for *inversely*, moving backwards instead of forwards. To fix this, we use our heading transformation function on our error. This lets us, in effect, rotate our PDFL controller along with our robot, increasing it's domain of validity, even for small rotation errors. We also remove the other linearity in our model, the control input function, by taking it's gradient around our predicted state and using that to adjust our PDFL, removing the nonlinearity as we're following the path.

5 Current/Future Work

While our team originally decided to use a differential drivetrain, after our first qualifier we realized that our choice did not result in the competitive edge we had hoped for. We wanted to be fast and resistant to other robots pushing our robot around, but found that head-on we had the same defensive power as other robots with more versatile drivetrains. So, we are instead opting for a traditional holonomic drivetrain—Mecanum drive. This uses four independently controlled actuators to enable three directions of movement - drive, strafing side to side, and turning.

I am actively in the process of porting our pathing system over to this new drivetrain. Apart from a slightly different dynamic model, the biggest change is to the control limiter function. We have control in drive, strafe, and heading, but these do not nicely map to motor powers—we have to find for parameters that fit a three parameter target control. The additional constraint we impose is that the wheels should be pushing against each other as little as possible. From this, we can generate a 4x3 wheel power matrix W that maps from a Drive-Strafe-Heading control vector (d, s, h) to (fl, fr, bl, br) wheel power vector.

What this means for the pathing is that we have to limit our wheel powers in this space, which we can do by changing the control limiting function to $s'(u) = W^L s(Wu)$. Associated derivatives are again calculated analytically.